LiveCodeBench-Pro: GPT-6 Astra Methodology

Same 10 problems as the DeepSeek V4 Flash run of 20 October 2026.
Same C++ judge: LightCPVerifier on `linux/arm64`.
Dataset revision: `adebffce047d`.
Upstream commit: `d614eef4f95b`.

GPT-6 Astra: solutions authored in Codex on 21 October 2026, with `xhigh` reasoning. All ten sources frozen before testing.
One judged submission per problem, with no revisions after feedback.
No editorials, reference solutions or hidden test inputs consulted.

Generation settings differ: DeepSeek used temperature 0 and a 4,096-token API budget. Astra used a shared Codex conversation without a per-problem token cap.
These are subset results under different generation conditions.
Source: LiveCodeBench-Pro

LiveCodeBench-Pro: GPT-6 Astra Results

Metric	GPT-6 Astra	DeepSeek V4 Flash
Accepted / selected	10/10	4/10
Acceptance	100.0%	40.0%
Completed judging	10/10	10/10
Failures	0	0
Compilation Error	0	4
Wrong Answer	0	2
Mean inference latency	Not measured	25.32 s

Acceptance = accepted / selected problems. Compilation errors and wrong answers count as completed judging, but do not count as accepted.
Astra's Codex workflow has no comparable per-request latency measurement.
Model IDs: `gpt-6-astra` and `deepseek-ai/DeepSeek-V4-Flash`.
This ten-problem subset is not a full-benchmark leaderboard score.

LiveCodeBench-Pro: Verdicts for All Ten Problems

ID	Challenge	GPT-6 Astra	DeepSeek
1983A	Array Divisibility	Accepted	Accepted
1983B	Corner Twist	Accepted	CE
1983C	Have Your Cake and Eat It Too	Accepted	Accepted
1983D	Swap Dilemma	Accepted	Accepted
1983E	I Love Balls	Accepted	WA
1983F	array-value	Accepted	Accepted
1983G	Your Loss	Accepted	CE
1988A	Split the Multiset	Accepted	CE
1988B	Make Majority	Accepted	CE
1988C	Increasing Sequence with Fixed OR	Accepted	WA

CE = Compilation Error. WA = Wrong Answer.
Astra passed 582/582 recorded judge cases across the ten submissions.
Verdicts use the same judge, dataset revision and problem IDs as DeepSeek.